

Mathematics of Deep Learning

-

Practical Session

Jan Macdonald & Rafael Reisenhofer

Mathematisches Forschungszentrum Oberwolfach
16 October 2018

Outline

1. Approximation of functions in Sobolev spaces with deep networks
2. TensorFlow Tutorial I (Proof of upper bound)
3. Approximation of functions in Sobolev spaces with shallow networks
4. TensorFlow Tutorial II (Experimental verification)
5. Discussion of Results

Some Definitions

- $\mathcal{W}^{r,p}([0,1]^d)$: Sobolev space of r -times weakly differentiable functions in $L^p([0,1]^d)$ with norm

$$\|f\|_{\mathcal{W}^{r,p}([0,1]^d)} = \begin{cases} \left(\sum_{|\alpha| \leq r} \|D^\alpha f\|_p^p \right)^{1/p} & 1 \leq p < \infty, \\ \max_{|\alpha| \leq r} \|D^\alpha f\|_\infty & p = \infty \end{cases}$$

- $F_{r,d}^p$: Unit ball in $\mathcal{W}^{r,p}([0,1]^d)$.
- Activation functions
 - **ReLU**: $\sigma(x) = \max(0, x)$
 - **Sigmoid**: $\sigma(x) = \frac{1}{1+e^{-x}}$

Approximation of Functions in $F_{r,d}^\infty$ with Deep Networks

Theorem – Upper Bound (Yarotsky, 2017)

*For any d, r , and $\epsilon \in (0, 1)$, there is a ReLU network architecture that can express any function from $F_{r,d}^\infty$ with error ϵ with depth at most $C(\ln(1/\epsilon) + 1)$ and **at most** $C\epsilon^{-d/r}(\ln(1/\epsilon) + 1)$ **weights**, with some constant C depending on d and r .*

Theorem – Lower Bound (Yarotsky, 2017)

*Fix d and r . For any $\epsilon \in (0, 1)$, a ReLU network architecture capable of approximating any function $f \in F_{r,d}^\infty$ with error ϵ must have **at least** $C\epsilon^{-d/2r}$ **weights**, with some constant C depending on d and r .*

Proof Idea for Upper Bound

The proof that a ReLU architecture with depth at most $C(\ln(1/\epsilon) + 1)$ and at most $C\epsilon^{-d/r}(\ln(1/\epsilon) + 1)$ connections suffices to obtain an approximation error ϵ for any function in $F_{r,d}^\infty$ requires the following steps:

1. $f(x) = x^2$ on $[0, 1]$ can be approximated with arbitrary precision with $O(\ln(1/\epsilon))$ layers (explicit construction via iterated sawtooth functions)

Proof Idea for Upper Bound

The proof that a ReLU architecture with depth at most $C(\ln(1/\epsilon) + 1)$ and at most $C\epsilon^{-d/r}(\ln(1/\epsilon) + 1)$ connections suffices to obtain an approximation error ϵ for any function in $F_{r,d}^\infty$ requires the following steps:

1. $f(x) = x^2$ on $[0, 1]$ can be approximated with arbitrary precision with $O(\ln(1/\epsilon))$ layers (explicit construction via **iterated sawtooth functions**)
2. Due to $xy = ((x + y)^2 - x^2 - y^2)/2$, this implies that multiplication can be efficiently implemented.

Proof Idea for Upper Bound

The proof that a ReLU architecture with depth at most $C(\ln(1/\epsilon) + 1)$ and at most $C\epsilon^{-d/r}(\ln(1/\epsilon) + 1)$ connections suffices to obtain an approximation error ϵ for any function in $F_{r,d}^\infty$ requires the following steps:

1. $f(x) = x^2$ on $[0, 1]$ can be approximated with arbitrary precision with $O(\ln(1/\epsilon))$ layers (explicit construction via **iterated sawtooth functions**)
2. Due to $xy = ((x+y)^2 - x^2 - y^2)/2$, this implies that multiplication can be efficiently implemented.
3. Approximate arbitrary f as

$$\tilde{f} = \sum_n \varphi_n P_n,$$

where φ_n are piecewise linear functions that define a **partition of unity** and P_n are **local Taylor polynomials**.

Proof Idea for Upper Bound

The proof that a ReLU architecture with depth at most $C(\ln(1/\epsilon) + 1)$ and at most $C\epsilon^{-d/r}(\ln(1/\epsilon) + 1)$ connections suffices to obtain an approximation error ϵ for any function in $F_{r,d}^\infty$ requires the following steps:

1. $f(x) = x^2$ on $[0, 1]$ can be approximated with arbitrary precision with $O(\ln(1/\epsilon))$ layers (explicit construction via **iterated sawtooth functions**)
2. Due to $xy = ((x+y)^2 - x^2 - y^2)/2$, this implies that multiplication can be efficiently implemented.
3. Approximate arbitrary f as

$$\tilde{f} = \sum_n \varphi_n P_n,$$

where φ_n are piecewise linear functions that define a **partition of unity** and P_n are **local Taylor polynomials**.

4. Construct a ReLU network that approximates φ_n , P_n , and the product $\varphi_n P_n$ sufficiently well.

TensorFlow Tutorial Part I

https://github.com/jmaces/mfo_dl_seminar_2018

Approximation of Functions in $F_{r,d}^p$ with Shallow Networks

Theorem – Upper Bound (Mhaskar, 1996)

Let $1 \leq p \leq \infty$. For any r, d , and $\epsilon \in (0, 1)$, there is a network with a **smooth activation function** and only **one hidden layer** that can express any function from $F_{r,d}^p$ with error ϵ **at most** $C\epsilon^{-d/r}$ **weights**, with some constant C depending on d and r .

Approximation of Functions in $F_{r,d}^p$ with Shallow Networks

Theorem – Upper Bound (Mhaskar, 1996)

*Let $1 \leq p \leq \infty$. For any r, d , and $\epsilon \in (0, 1)$, there is a network with a **smooth activation function** and only **one hidden layer** that can express any function from $F_{r,d}^p$ with error ϵ **at most** $C\epsilon^{-d/r}$ **weights**, with some constant C depending on d and r .*

Observation: Deep ReLU networks and shallow networks with smooth activation functions seem to have a similar approximation behavior with respect to functions in $F_{r,d}^p$.

Approximation of Functions in $F_{r,d}^p$ with Shallow Networks

Theorem – Upper Bound (Mhaskar, 1996)

Let $1 \leq p \leq \infty$. For any r, d , and $\epsilon \in (0, 1)$, there is a network with a **smooth activation function** and only **one hidden layer** that can express any function from $F_{r,d}^p$ with error ϵ **at most** $C\epsilon^{-d/r}$ **weights**, with some constant C depending on d and r .

Observation: Deep ReLU networks and shallow networks with smooth activation functions seem to have a similar approximation behavior with respect to functions in $F_{r,d}^p$.

Can we verify this numerically?

Experimental Setup – Network Topologies

- We consider network architectures with different parameters in three dimensions:

Activation Function:	$\{\text{ReLU, sigmoid}\}$
Depth:	$\{1, \dots, 18\}$ (hidden layers)
Width:	$\{5, 10, \dots, 50\}$

- We allow interconnections between layers: Each computational node (neuron) is connected to **all nodes from all earlier layers**.
- In the cases of only 1, 2, and 3 hidden layers, we consider widths up to 500.

Experimental Setup – Training and Testing

- We try to learn the sum of a smooth function $B(x, y)$ and a function $g_{r_1, r_2}(x, y)$ in $\mathcal{W}^{\min\{r_1, r_2\}, p}([0, 1]^d)$ on $[0, 1]^2$:

$$B(x, y) = \underbrace{\binom{n_1}{k_1}_1 x^{n_1} (1-x)^{n_1-k_1} \binom{n_2}{k_2} y^{n_2} (1-y)^{n_2-k_2}}_{\text{smooth 2D Bernstein polynomial}},$$

$$g_{r_1, r_2}(x, y) = \underbrace{\text{sgn}(x - 0.4)(x - 0.4)^{r_1} \text{sgn}(y - 0.6)(y - 0.6)^{r_2}}_{\text{only } \min\{r_1, r_2\} \text{ weak derivatives at } (0.4, 0.6)},$$

$$f(x, y) = B(x, y) + g_{r_1, r_2}(x, y)$$

Experimental Setup – Training and Testing

- ▶ We try to learn the sum of a smooth function $B(x, y)$ and a function $g_{r_1, r_2}(x, y)$ in $\mathcal{W}^{\min\{r_1, r_2\}, p}([0, 1]^d)$ on $[0, 1]^2$:

$$B(x, y) = \underbrace{\binom{n_1}{k_1}_1 x^{n_1} (1-x)^{n_1-k_1} \binom{n_2}{k_2} y^{n_2} (1-y)^{n_2-k_2}}_{\text{smooth 2D Bernstein polynomial}},$$

$$g_{r_1, r_2}(x, y) = \underbrace{\text{sgn}(x - 0.4)(x - 0.4)^{r_1} \text{sgn}(y - 0.6)(y - 0.6)^{r_2}}_{\text{only } \min\{r_1, r_2\} \text{ weak derivatives at } (0.4, 0.6)},$$


$$f(x, y) = B(x, y) + g_{r_1, r_2}(x, y)$$

- ▶ Training: 100k iterations of gradient descent with uniformly sampled training points and batch size 4096.
- ▶ Loss function: $\|\cdot\|_2$
- ▶ Learning rate: Exponential decrease from 0.2 to 0.001.
- ▶ Evaluation: $\|\cdot\|_\infty$ on 200×200 equidistant grid.

TensorFlow Tutorial II


Results – Best ReLU vs. Best Sigmoid

► Best ReLU-based approximation




depth: 11, width: 35, $M = 44731$, ℓ_∞ -error ≈ 0.00372 .

► Best sigmoid-based approximation

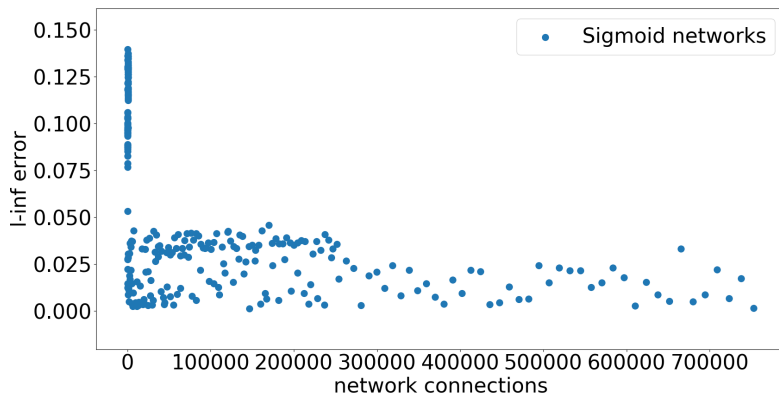


depth: 5, width: 220, $M = 146521$, ℓ_∞ -error ≈ 0.00013


Size vs. Error in ReLU Networks



Size vs. Error in Shallow Sigmoid Networks



Size vs. Error in 3-Layer Sigmoid Networks



References

- ▶ D. Yarotsky. “Error bounds for approximations with deep ReLU networks”. In: **Neural Networks** 94 (2017), pp. 103–114
- ▶ H. N. Mhaskar. “Neural networks for optimal approximation of smooth and analytic functions”. In: **Neural computation** 8.1 (1996), pp. 164–177